

# Lecture 1b/2a: Introduction to PINNS

**Chris Budd and Aengus Roberts<sup>1</sup>**

<sup>1</sup>University of Bath

CWI, October 2025

## Some papers to look at

- Raissi, M., P. Perdikaris, and G. E. Karniadakis. 2019. “Physics-Informed Neural Networks: A Deep Learning Framework for Solving Forward and Inverse Problems Involving Nonlinear Partial Differential Equations.” *Journal of Computational Physics* 378: 686–707. <https://doi.org/https://doi.org/10.1016/j.jcp.2018.10.045>.
- Grossman et. al. *Can PINNS beat the FE method?*
- Shin et. al. *On the convergence of PINNS for linear elliptic and parabolic PDEs*
- Wang et. al. *When and why PINNS fail to train: a neural tangent kernel perspective*
- Kiyani *Which Optimizer Works Best for Physics-Informed Neural Networks and Kolmogorov-Arnold Networks?*
- Chen et. al. *Sidecar: A Structure-Preserving Framework for Solving Partial Differential Equations with Neural Networks*
- Russell et. al. *Two-point boundary value problems*

# Motivation: Solving ODEs and PDEs

Seek to solve ODE/PDE problems of the form

$$\mathbf{u}_t = F(\mathbf{x}, u, \nabla u, \nabla^2 u) \quad \text{with BC}$$

eg. ODE (IVP or BVP)

$$\frac{du}{dt} = u^2, u(0) = 1, \quad -\epsilon^2 \frac{d^2 u}{dx^2} + u = 1, \quad u(0) = u(1) = 0.$$

eg. PDE

$$-\Delta u = f(x), \quad iu_t + \Delta u + u|u|^2 = 0, \quad u_t = \Delta u + f(x, u) \quad x \in R^n$$

# Classical Method 1. Finite Differences

Work with a set of **point values** for a function **never with a function directly**  
IVP/BVP:

$$U_j^n \approx u(n\Delta t, j\Delta x), \quad \Delta t, \Delta x \ll 1$$

eg.

$$u_t = u_{xx} + u^3$$

IMEX Crank-Nicholson method:

$$\frac{U_j^{n+1} - U_j^n}{\Delta t} = \frac{1}{2\Delta x^2} \left( U_{j+1}^{n+1} - 2U_j^{n+1} + U_{j-1}^{n+1} + U_{j+1}^n - 2U_j^n + U_{j-1}^n \right) + (U_j^n)^3.$$

- Have error estimates of the form

$$\|U_j^n - u(n\Delta t, j\Delta x)\| < C(u)\Delta t^p \Delta x^q, \quad \Delta t, \Delta x \rightarrow 0$$

- Can reduce  $C(u)$  and increase  $p, q$  using an adaptive approach.
- Lots of software
- Have to recover the function from the point values
- Awkward in higher dimensions!

## Classical Method 2: Finite Elements

Express  $u(x, t)$  as a Galerkin approximation:

$$u(t, x) \approx U(t, x) = \sum_{i=0}^N U_i(t) \phi_i(x)$$

with  $\phi_i(x)$  **Not** globally differentiable *locally supported, piece-wise polynomial spline functions*

- Work with a function  $U(x) \in H^1$
- Require  $U$  to satisfy a *weak form* of the PDE (see Lecture 4)
- Have guaranteed error estimates of the form

$$\|u - U\|_{H^1} < C(u)N^{-\alpha}$$

- Lots of software (see Lecture 2b)
- Can reduce  $C(u)$  and increase  $\alpha$  using an adaptive approach.
- Awkward in higher dimensions!

## Classical Method 3: Collocation

Express  $u(x, t)$  as a function approximation:

$$u(t, x) \approx U(t, x) = \sum_{i=0}^N U_i(t) \phi_i(x)$$

with  $\phi_i(x)$  :

**Locally differentiable locally supported, piece-wise polynomial spline functions such as Lagrange polynomial interpolants**

- Work with a function  $U(x) \in C_{loc}^n$
- Require  $U$  to satisfy the PDE **exactly** at a carefully chosen set of **collocation points**

- Have guaranteed error estimates of the form

$$\|u - U\|_{C_2} < C(u)N^{-\alpha}$$

- Can reduce  $C(u)$  and increase  $\alpha$  using VERY carefully chosen **collocation points**
- Implemented in `python` as `solvebvp`, `Colsys`
- Impossible in higher dimensions!



# Issues with classical methods

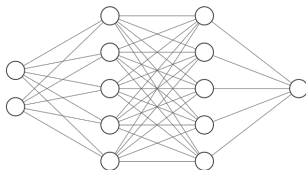
- Well tested and good convergence theories
- Accuracy of the computation depends crucially on the choice and shape of the mesh points
- Generally require **time and practice** to implement (although FireDrake etc. makes this easier)
- Two-fold curse of dimensionality

# PINNS

Physics Informed Neural Networks for solving PDEs: advertised as "Mesh free methods".

Use a Deep Neural Net of width  $W$  and depth  $L$  to give a **nonlinear functional approximation** to  $u(\mathbf{x})$  with **input**  $x$ .

$$y(\mathbf{x}) = DNN(\mathbf{x})$$



$y(\mathbf{x})$  is constructed via a combination of linear transformations and nonlinear/semi-linear activation functions.

Example: Shallow 1D neural net

$$y(\mathbf{x}) = \sum_{i=0}^{W-1} c_i \sigma(a_i \mathbf{x} + \mathbf{b}_i)$$

Often take

$$\sigma(z) = \text{ReLU}^3(z) \equiv z_+^3, \quad \text{or} \quad \sigma(z) = \tanh(z)$$

As a result we can define  $y_{zz}$  for every point  $z$

# Operation of a 'traditional' PINN

Want to solve a ODE/PDE in  $\mathbf{x}$

- PINNS are similar to collocation methods
- Assume that  $y(\mathbf{x})$  has strong regularity eg.  $C^2$
- Differentiate  $y(\mathbf{x})$  **exactly** using the chain rule
- Evaluate the **PDE residual** at  $N_r$  **collocation points**  $\mathbf{X}_i$ , :chosen to be uniformly spaced, or **samples from a random distribution**
- Train the neural net to minimise a **loss function**  $L$  combining the PDE residual and boundary and initial conditions. May use **Epochs linked to the random training samples**.

## Eg 1. Solution of regular two-point BVPs by PINNs

Consider the two-point BVP with Dirichlet boundary conditions:

$$-u_{xx} = f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b.$$

Define output of the PINN by  $y(x)$  and residual  $r(x) := y_{xx} + f(x, y, y_x)$ .  
Train the coefficients  $\theta$  of the NN by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y(1) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** (possibly randomly) placed in  $(0, 1)$ .

## Eg 2. Solution of regular second order IVPs by PINNs

Consider the second order IVP with initial conditions:

$$-u_{xx} = f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u'(0) = b.$$

Define output of the PINN by  $y(x)$  and residual  $r(x) := y_{xx} + f(x, y, y_x)$ .  
Train the coefficients  $\theta$  of the NN by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y'(0) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** (possibly randomly) placed in  $(0, 1)$ .

## Eg 3. Solution of regular parabolic PDEs by time-stepping PINNs

Consider the **semilinear parabolic PDE** with Dirichlet boundary conditions:

$$u_t = u_{xx} + f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b$$

and initial condition

$$u(x, 0) = u_0(x).$$

**Implicit time-stepping scheme**  $U^n(x) \approx u(n\Delta t, x)$ .

$$\frac{U^{n+1}(x) - U^n(x)}{\Delta t} = U_{xx}^{n+1} + f(x, U^{n+1}, U_x^{n+1}) \equiv F(U^{n+1}).$$

- Start with  $U^0(x) = u_0(x)$
- For  $n > 0$  : Define output of the PINN by  $y(x)$  and residual

$$r(x) := U^n(x) + \Delta t F(y).$$

- The PINN is trained by minimising the loss function

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(X_i^r)|^2 + \frac{1}{2} (|y(0) - a|^2 + |y(1) - b|^2),$$

where  $\{X_i^r\}_i^{N_r}$  are the **collocation points** placed in  $(0, 1)$ .

- Set  $U^{n+1} = y(x)$  and repeat.



## Eg 4. Solution of regular parabolic PDEs by a full PINN

Consider the same semilinear parabolic PDE

$$u_t = u_{xx} + f(x, u, u_x), \quad x \in [0, 1] \quad u(0) = a, \quad u(1) = b.$$

**Bold approach!**

- Have NN with  $y(t, x)$  a function of  $t$  and  $x$
- Take  $N_r$  collocation points  $Z_i$  in **space and time**
- Residual  $r = u_t - u_{xx} - f(x, u, u_x)$
- Minimise

$$L = \frac{1}{N_r} \sum_i^{N_r} |r(Z_i^r)|^2 + \frac{\beta}{2} (|y(0) - a|^2 + |y(1) - b|^2) \\ + \gamma \|y(0, x) - u_0(x)\|^2.$$

**Problematic as space and time play very different roles in the PDE**

Numerical results for:  $-u'' = \pi^2 \sin(\pi x)$ .

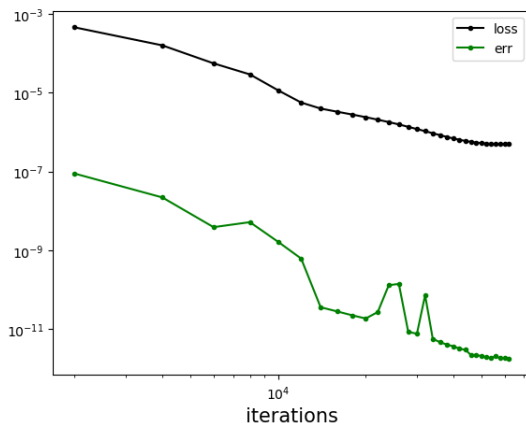
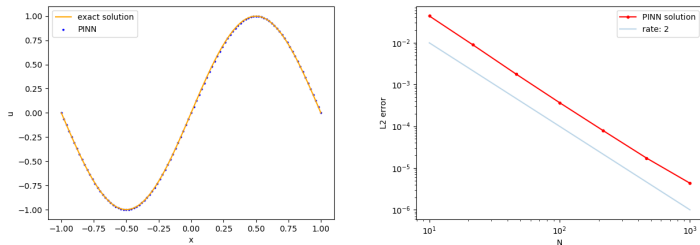


Figure: Residual based Loss and  $L^2$  error of the PINN solution for  $N_r = 100$


# Numerical results: $u(x) = \sin(\pi x)$



**Figure:** (Left) PINN with depth  $L = 2$  and width  $W = 30$ .  $N_r = 100$  uniformly distributed quadrature points, activation function:  $\tanh$ , optimizer: Adam with  $lr = 1e - 3$  (Right) Convergence rate for 1st order interpolant

## Eg 2. Singular Reaction-Diffusion Equation

Solve  $-\varepsilon^2 u_{xx} + u = 1 - x$  on  $[0, 1]$   $u(0) = u(1) = 0$



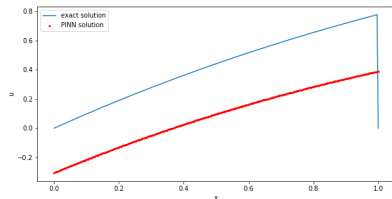
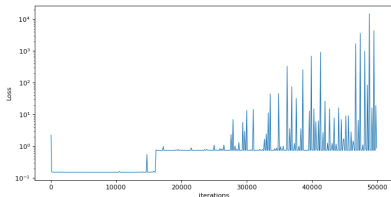
**Figure:** PINN (tanh) trained for 20000 epochs,  $N_r = 101$ , Adam optimizer with  $lr = 1e - 3$ . (left) Uniform collocation points (right) Adapted collocation points  
**much faster training**

## Eg 3. Bad news: Convection-dominated equation

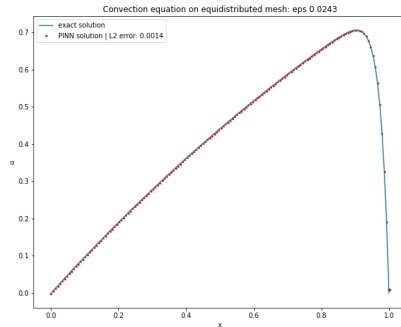
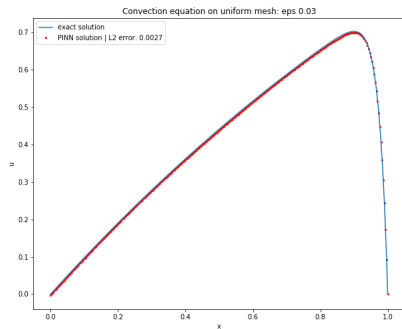
PINNs often fail to train when the solution of the BVP exhibits singular and convective behaviour [Krishnapriyan, Aditi et. al., (2021)]:

$$-\varepsilon u_{xx} + \left(1 - \frac{\varepsilon}{2}\right) u_x + \frac{1}{4} \left(1 - \frac{1}{4}\varepsilon\right) u = e^{-x/4} \text{ on } [0, 1] \quad u(0) = u(1) = 0$$

$$u(x) = \exp^{\frac{-x}{4}} \left( x - \frac{\exp^{-\frac{1-x}{\varepsilon}} - \exp^{-\frac{1}{\varepsilon}}}{1 - \exp^{-\frac{1}{\varepsilon}}} \right)$$



# Numerical results: Convection Equation



# General questions for consideration

A PINN is based on a NN so we can expect in theory to achieve the expressivity and highly accurate approximation results. But do we see this in practice?

- ① When do and don't PINNs work, and why?
- ② How do these answers depend on (i) the problem (ii) choice of activation function, optimisation, collocation points, conditioning etc
- ③ Can we develop a useful convergence theory for a PINN using tools from approximations theory, bifurcation theory, numerical analysis etc.
- ④ How does a PINN compare to a finite element method?

# Convergence theory for PINNS

[Shin, Darbon and Kaniardarkis, 2020], [Jiao, Lai, Lo, Wang, Yang, 2024]

- PINN error is a combination of approximation error and training/optimization error
- Show that a PINN (depth  $L$  width  $W$ ) can be constructed with low approximation error which reduces as the complexity of the PINN increases.
- Hope that the optimization error can be reduced to acceptable levels.
- Try things out on simple problems



## BUT

- **Non convex (no guarantee of uniqueness or convergence of training)**
- PINNs are nonlinear function approximators. **No equivalent of Cea's lemma giving a bound on the solution error.**
- Solutions can sometimes have no connection to reality!

# Simple convergence theory for any collocation method including a PINN

Consider the DE:

$$-u_{xx} = f(x), \quad u(0) = f, u(1) = g.$$

There exists an exact solution using a Green's function  $G(x, y)$

$$u(x) = G * f = \int_0^1 G(x, y) f(y) dy.$$

Suppose we have  $N$  collocation points  $Y_i$  then we can approximate the integral by a quadrature

$$u(x) = \sum_{i=1}^N w_i G(x, Y_i) f(Y_i) + \mathcal{O}(N^{-\alpha})$$

for some  $\alpha$  depending on the weights  $w_i$  and the smoothness of  $u$ .

Now let  $U(x)$  be any solution of the collocation problem

$$-U_{xx}(Y_i) = f(Y_i).$$

It follows that

$$\sum_{i=1}^N w_i G(x, Y_i) f(Y_i) = - \sum_{i=1}^N w_i G(x, Y_i) U_{xx}(Y_i) = U(x) + \mathcal{O}(N^{-\beta})$$

for some  $\beta$  depending on the weights  $w_i$  and the smoothness of  $u$ .  
Hence we have the convergence result

$$u(x) = U(x) + \mathcal{O}(N^{-\alpha}) + \mathcal{O}(N^{-\beta}).$$

Can do better with adaptivity and careful choice of collocation points: see [Russell et. al.]

# Implementation of collocation methods and PINNS

A standard collocation method for the problem  $-u_{xx} = f(x)$  is a linear approximation of the form:

$$u(x) \approx U(x) = \sum a_j \phi_j(x).$$

The collocation conditions at the points  $X_i$  lead to a linear equation:

$$\sum_j A_{i,j} a_j = f_i$$

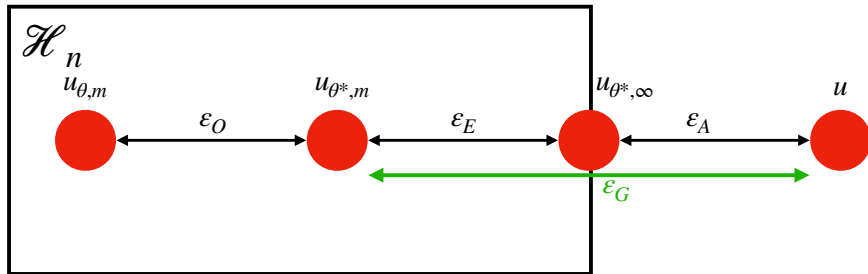
$$A_{i,j} = -\epsilon^2 \phi_j(X_i) + \phi_j(X_i), \quad f_i = f(X_i).$$

These linear equations are (in general) well conditioned and have a unique solution

Compare with the PINN formulation for the same problem

- PINN is a nonlinear approximation
- It may have better expressivity
- The collocation equations become an **ill conditioned** optimisation problem
- We may not find a good optimum due to the **lack of convexity and the ill conditioning**
- Pre conditioning and the right choice of optimiser are essential

# Detailed PINN Convergence Theory 1



# Detailed PINN Convergence Theory 2

Underlying solution:  $u(x)$

$H_n$ : Class of functions approximated by NN with  $n$  degrees of freedom

- $u_{\theta,m}$ : Approximation found in practice after some training. Is less accurate than
- $u_{\theta_m}^*$ : Approximation obtained by perfect optimisation with finite collocation points. Is less accurate than
- $u_{\theta^*,\infty}$ : Approximation obtained by perfect optimisation with infinite collocation points.

Measured error is  $\|u_{\theta,m} - u\|$

- $\|u_{\theta,m} - u_{\theta^*,m}\|$ : Optimisation error  $\mathcal{E}_O$ : Hard to control and depends on the optimiser and the initialisation: See Lecture 1a
- $\|u_{\theta^*,m} - u_{\theta^*,\infty}\|$ : Estimation error  $\mathcal{E}_E$ : Main results on this
- $\|u_{\theta^*,\infty} - u^*\|$ : Approximation error  $\mathcal{E}_A$ : See Lecture 1a



# Training data

[Shin et. al] prove

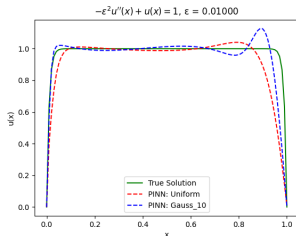
**Theorem** *If there is a random set of training data and  $n$  samples are chosen for training. Then the **optimal minimiser**  $h_n$  over this set converges to the best approximation  $\hat{h}$  as  $n \rightarrow \infty$ .*

# Results on a mildly singular BVP

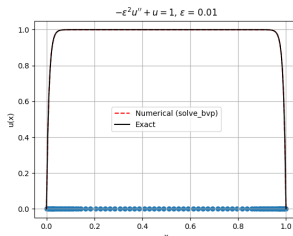
Compare standard collocation and a PINN for the ODE using ADAM

$$-\epsilon^2 u_{xx} + u = 1, \quad u(0) = u(1) = 0, \quad 0 < \epsilon \ll 1$$

This has a boundary layer at  $x = 0, x = 1$  of width  $\epsilon$ .

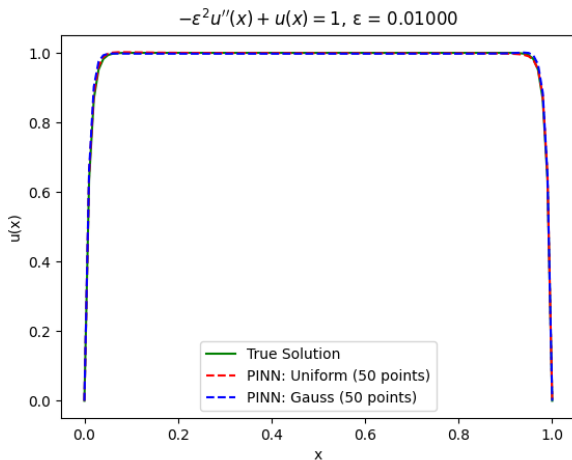


PINN



Collocation

# Using the LBFGS optimiser



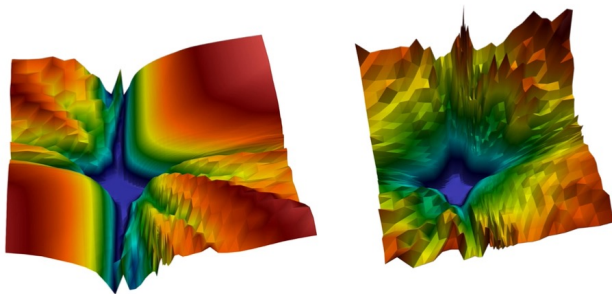
# How and why do PINNS struggle to train .. and how can we fix this?

A subject of intensive research! See [Wang et. al.] for a summary:  
Observe

- PINNS train especially badly when they have high frequency features
- Loss function has multi-scale interactions, leading to ill-conditioning and stiffness in the gradient flow dynamics and a stringent stability requirement on the loss rate
- Standard networks show ill-conditioning and spectral bias and cannot learn high frequencies. This has been seen for example in weather forecasting

They advocate the use of Neural Tanjent Kernels in the optimisation process. [Kiyani et. al.] extend this with the development of better optimisers eg. Broyden methods, for PINNS

Loss landscape is highly non-convex, [Kiyani et. al.]



**Also** Issues with **conditioning** as in the last lecture on approximation.  
Can fix for shallow networks by **preconditioning**